



# Statistical Machine Learning

## *Lecture 2 – Linear regression, regularization*



UPPSALA  
UNIVERSITET

David Sumpter

Division of Systems and Control

Department of Information Technology

Uppsala University.

# Summary of Lecture 1 (I/III)

---

What is this course about? **Supervised** machine learning

In one sentence:

Methods for automatically learning (training, estimating, ...)

**a model** for the relationship between

- the **input**  $\mathbf{x}$ , and
- the **output**  $y$

from observed **training data**

$$\mathcal{T} := \{(y_1, \mathbf{x}_1), (y_2, \mathbf{x}_2), \dots, (y_n, \mathbf{x}_n)\}.$$

# Summary of Lecture 1 (II/III)

---

## Regression vs. classification

- **Numerical** variables take on numerical values (real numbers, integer values, . . . ).
- **Categorical** variables take on values in one of  $K$  distinct classes, e.g. “true or false”, “disease type  $A$ ,  $B$  or  $C$ ”.

**Regression** is when the output  $y$  is numerical.

**Classification** is when the output  $y$  is categorical.

# Summary of Lecture 1 (II/III)

---

What maths do we need.

- **Calculus** Finding a parameter which minimizes the distance between two points.
- **Matrix algebra** For keeping track of sum of squares.
- **Probability theory** Estimating parameters. Normal distribution important

One key idea that brings these together is maximum likelihood.  
Finding the model that is maximally likely (closest to data).



# Where are we now?

---

# Outline – Lecture 2

---

**Aim:** To introduce linear regression and its regularized version.

## Outline:

1. Summary of Lecture 1
2. Linear regression models
3. Maximum likelihood and least squares
4. Regularization
  - Ridge regression
  - LASSO

---

*Linear regression is the foundation of statistics and (supervised) machine learning.*



# The course book

---

We have developed this course over the last 4 years.

A book based on these notes will soon be published as a textbook with Cambridge University Press.

**Please read in parallel to your studies.**

Book website:

[smlbook.org](https://smlbook.org)

Help us by reporting errors via practical GitHub interface:

[github.com/uu-sml/sml-book-page/issues](https://github.com/uu-sml/sml-book-page/issues)

# Regression



- **Input variable**  $X$
- **Output variable**  $Y$

**Regression:** learning a model explaining  $Y$  from  $X$ , when  $Y$  is numerical.

$$Y = f(X; \beta) + \epsilon$$

$\beta$  are the **parameters** of the model

( $Y$  categorical  $\rightarrow$  classification)





# Numerical or categorical variables

---

## Numerical or categorical?

17.31 kg, 22.37 kg, 51.34 kg

1 = brown hair, 2 = red hair, 3 = blonde hair

Adenine, Thymine, Cytosine, Guanine

1 bike, 2 bikes, 5 bikes

# Numerical or categorical variables

---

## Numerical or categorical?

17.31 kg, 22.37 kg, 51.34 kg

Numerical

1 = brown hair, 2 = red hair, 3 = blonde hair

Adenine, Thymine, Cytosine, Guanine

1 bike, 2 bikes, 5 bikes

# Numerical or categorical variables

---

## Numerical or categorical?

17.31 kg, 22.37 kg, 51.34 kg

Numerical

1 = brown hair, 2 = red hair, 3 = blonde hair

Categorical

Adenine, Thymine, Cytosine, Guanine

1 bike, 2 bikes, 5 bikes

# Numerical or categorical variables

---

## Numerical or categorical?

17.31 kg, 22.37 kg, 51.34 kg

Numerical

1 = brown hair, 2 = red hair, 3 = blonde hair

Categorical

Adenine, Thymine, Cytosine, Guanine

Categorical

1 bike, 2 bikes, 5 bikes

# Numerical or categorical variables

---

## Numerical or categorical?

17.31 kg, 22.37 kg, 51.34 kg

Numerical

1 = brown hair, 2 = red hair, 3 = blonde hair

Categorical

Adenine, Thymine, Cytosine, Guanine

Categorical

1 bike, 2 bikes, 5 bikes

Numerical

# Numerical or categorical variables

---


## Numerical or categorical?

17.31 kg, 22.37 kg, 51.34 kg	Numerical
1 = brown hair, 2 = red hair, 3 = blonde hair	Categorical
Adenine, Thymine, Cytosine, Guanine	Categorical
1 bike, 2 bikes, 5 bikes	Numerical


Categorical output variable? → classification. (But can be done with regression!)

Categorical input variable? Still regression!

# Regression example: car stopping distances




# Regression example: car stopping distances





# Regression example: car stopping distances



*(in fact a linear regression model with nonlinear transformation of the input variables)*

# Regression example: Alpha Go zero

---



# Regression example: Alpha Go zero



- Input: State of the game ( $19 \times 19$  grid, either black, white or blank)
  - Output: Probability for the current player to win the game
- + *reinforcement learning*

Silver et al. **Mastering the game of Go with deep neural networks and tree search**, *Nature* 529, 484–489, 2016.



# Regression example: Alpha Go zero



- Input: State of the game ( $19 \times 19$  grid, either black, white or blank)
- Output: Probability for the current player to win the game

+ *reinforcement learning*

Silver et al. **Mastering the game of Go with deep neural networks and tree search**, *Nature* 529, 484–489, 2016.

- Input: Same
- Output: Probability for the current player to win the game *and* what move to make

Silver et al. **Mastering the game of Go without human knowledge**, *Nature* 550, 354–359, 2017.

Silver et al. **A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play**, *Science*, 362(6419): 1140–1144, 2018.

# Linear regression



*"Linear regression = Regression with a linear model",*

Output  $Y$  is linear combination of  $k$  inputs  $X_1, \dots, X_k$  plus some noise/error  $\epsilon$ ,

$$Y = \underbrace{\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k}_{f(X; \beta)} + \epsilon.$$

*Workflow (for most methods, not only linear regression):*


1. Learn/train/estimate model from training data  $\mathcal{T}$ :

find  $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_k$


2. Predict output for new test input using the model

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2 + \dots + \hat{\beta}_k X_k$$

# Linear regression ( $k = 1$ )



# Linear regression ( $k = 1$ )



# Learning the model from data



Linear regression model:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_k X_k + \epsilon$$

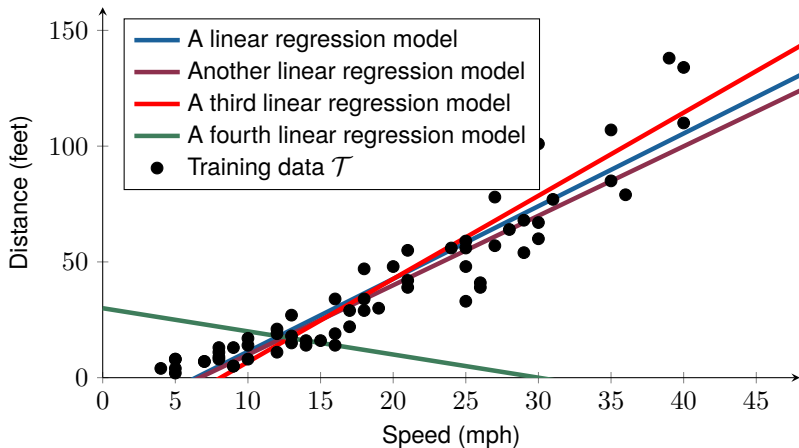
How to choose  $\beta_0, \beta_1, \dots, \beta_k$  ( $=\beta$ , column vector)?

Use **training data**  $\mathcal{T} = \{(y_i, x_i)\}_{i=1}^n$ !

$$x_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{ip} \end{bmatrix}, \quad \mathbf{X} = \begin{bmatrix} 1 & -x_1^T - \\ 1 & -x_2^T - \\ \vdots & \vdots \\ 1 & -x_n^T - \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}.$$




# What is a good model?






# Learning using maximum likelihood

Learning a model from data is a matter of looking at the errors  $\varepsilon$ !







# Learning using maximum likelihood

Learning a model from data is a matter of looking at the errors  $\varepsilon$ !







# Learning using maximum likelihood

Learning a model from data is a matter of looking at the errors  $\varepsilon$ !



# Learning using maximum likelihood

Learning a model from data is a matter of looking at the errors  $\varepsilon$ !



**Maximum likelihood:** Think of  $\varepsilon$  (dotted) as random variables, and *choose the model* (solid) *such that the resulting  $\varepsilon$  are as likely as possible.*

# Linear regression model in matrix form

---

Recall our linear regression model:

$$\mathbf{y} = \mathbf{X}\beta + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2 I).$$

Assumptions (for now):

1.  $\mathbf{y}$  – observed **random** variable.
2.  $\beta$  – unknown **deterministic** variable.
3.  $\mathbf{X}$  – known **deterministic** variable.
4.  $\varepsilon$  – unknown **random** variable.
5.  $\sigma_\varepsilon$  – unknown/known **deterministic** variable.

# Learning using maximum likelihood



Using the **maximum likelihood principle**

$$\hat{\beta} = \operatorname{argmax}_{\beta} P(\mathbf{y} | \mathbf{X}; \beta)$$

and assuming  $\epsilon \sim \mathcal{N}(0, \sigma_{\epsilon}^2)$  independently for each data point  $i$

$$\Rightarrow P(y_i | x_i; \beta) = \frac{1}{\sqrt{2\pi\sigma_{\epsilon}^2}} \exp \left( -\frac{1}{2\sigma_{\epsilon}^2} (\beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ip} - y_i)^2 \right)$$

$$\Rightarrow P(\mathbf{y} | \mathbf{X}; \beta) = \prod_{i=1}^n P(y_i | x_i; \beta) \propto \exp \left( -\frac{1}{2\sigma_{\epsilon}^2} \sum_{i=1}^n (\beta_0 + \dots + \beta_k x_{ip} - y_i)^2 \right)$$

$$\Rightarrow \hat{\beta} = \operatorname{argmin}_{\beta} \sum_{i=1}^n (\beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ip} - y_i)^2 = \operatorname{argmin}_{\beta} \underbrace{\|\mathbf{X}\beta - \mathbf{y}\|_2^2}_{\substack{\text{Loss function} \\ \text{induced by} \\ \text{maximum likelihood}}},$$

the **least squares** problem is achieved.

# Least squares in matrix form

---

The least squares problem

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{X}\beta - \mathbf{y}\|_2^2$$



# Least squares in matrix form

---

The least squares problem

$$V(\beta) = \|\mathbf{X}\beta - \mathbf{y}\|_2^2 = \beta^\top \mathbf{X}^\top \mathbf{X} \beta - 2\mathbf{y}^\top \mathbf{X} \beta + \mathbf{y}^\top \mathbf{y}$$

Minimize by differentiating and setting

$$\frac{\partial V(\beta)}{\partial \beta} = 2\mathbf{X}^\top \mathbf{X} \beta - 2\mathbf{X}^\top \mathbf{y}$$

Therefore,

$$\mathbf{X}^\top \mathbf{X} \hat{\beta} = \mathbf{X}^\top \mathbf{y}$$

# Least squares in matrix form

---

The least squares problem

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{X}\beta - \mathbf{y}\|_2^2$$

is solved by the **normal equations**

$$\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}.$$

Remember (from lecture 1)  $\mathbf{X}^\top \mathbf{X}$  is like sum of squares (similar to co-variances of input variables) and  $\mathbf{X}^\top \mathbf{y}$  is similar to co-variance between input and output.

For  $k = 1$ :

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

# Linear regression: the key concepts

---

**The linear regression model**

$$y = \beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k + \varepsilon$$

+


**Maximum likelihood**

$$\varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2) \text{ iid}$$

**Our first  
learning tool**

# Example


- $x$  = Speed
- $y$  = Distance



# Example

- $x$  = Speed
- $y$  = Distance

$$y = \beta_0 + \beta_1 x + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$$




# Example

- $x$  = Speed
- $y$  = Distance

$$y = \beta_0 + \beta_1 x + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$$

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

$$\mathbf{X} = \begin{bmatrix} 1 & 4 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 7 \\ 1 & 7 \\ 1 & 8 \\ \vdots & \vdots \\ 1 & 39 \\ 1 & 39 \\ 1 & 40 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 4 \\ 2 \\ 4 \\ 8 \\ 8 \\ 7 \\ 7 \\ 8 \\ \vdots \\ 138 \\ 110 \\ 134 \end{bmatrix}$$



# Example


- $x$  = Speed
- $y$  = Distance

$$y = \beta_0 + \beta_1 x + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$$

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

$$\mathbf{X} = \begin{bmatrix} 1 & 4 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 7 \\ 1 & 7 \\ 1 & 8 \\ \vdots & \vdots \\ 1 & 39 \\ 1 & 39 \\ 1 & 40 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 4 \\ 2 \\ 4 \\ 8 \\ 8 \\ 7 \\ 7 \\ 8 \\ \vdots \\ 138 \\ 110 \\ 134 \end{bmatrix}$$

$$\text{The normal equations} \Rightarrow \hat{\beta} = \begin{bmatrix} -20.1 \\ 3.1 \end{bmatrix}$$



# Example


- $x$  = Speed
- $y$  = Distance

$$y = \beta_0 + \beta_1 x + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$$

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

$$\mathbf{X} = \begin{bmatrix} 1 & 4 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 7 \\ 1 & 7 \\ 1 & 8 \\ \vdots & \vdots \\ 1 & 39 \\ 1 & 39 \\ 1 & 40 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 4 \\ 2 \\ 4 \\ 8 \\ 8 \\ 7 \\ 7 \\ 8 \\ \vdots \\ 138 \\ 110 \\ 134 \end{bmatrix}$$

$$\text{The normal equations} \Rightarrow \hat{\beta} = \begin{bmatrix} -20.1 \\ 3.1 \end{bmatrix}$$





# Example


- $x$  = Speed
- $y$  = Distance

$$y = \beta_0 + \beta_1 x + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$$

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

$$\mathbf{X} = \begin{bmatrix} 1 & 4 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 7 \\ 1 & 7 \\ 1 & 8 \\ \vdots & \vdots \\ 1 & 39 \\ 1 & 39 \\ 1 & 40 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 4 \\ 2 \\ 4 \\ 8 \\ 8 \\ 7 \\ 7 \\ 8 \\ \vdots \\ 138 \\ 110 \\ 134 \end{bmatrix}$$

$$\text{The normal equations} \Rightarrow \hat{\beta} = \begin{bmatrix} -20.1 \\ 3.1 \end{bmatrix}$$



Use the model for predictions!

# Transforming the inputs

---

*"If the speed  $x$  is an input variable, why can't the kinetic energy ( $\propto x^2$ ) be an input variable?"*

**We can make arbitrary nonlinear transformations to the input variables!**

The model is still a linear regression model, since

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 \cos(x) + \beta_4 \arctan(x) + \varepsilon$$

is equivalent to

$$\begin{aligned} y &= \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \varepsilon, \\ \text{with } X_1 &= v \\ X_2 &= v^2 \\ X_3 &= \cos(v) \\ X_4 &= \arctan(v) \end{aligned}$$

$x$  = original input variable,  $x_i$  transformed input variables (features).

# Example


- $x$  = Speed
- $y$  = Distance

$$y = \beta_0 + \beta_1 x + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$$

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

$$\mathbf{X} = \begin{bmatrix} 1 & 4 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 5 \\ 1 & 7 \\ 1 & 7 \\ 1 & 8 \\ \vdots & \vdots \\ 1 & 39 \\ 1 & 39 \\ 1 & 40 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 4 \\ 2 \\ 4 \\ 8 \\ 8 \\ 7 \\ 7 \\ 8 \\ \vdots \\ 138 \\ 110 \\ 134 \end{bmatrix}$$

The normal equations  $\Rightarrow \hat{\beta} = \begin{bmatrix} -20.1 \\ 3.1 \end{bmatrix}$



# Example


- $x$  = Speed
- $y$  = Distance

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$$

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix}$$

$$\mathbf{X} = \begin{bmatrix} 1 & 4 & 16 \\ 1 & 5 & 25 \\ 1 & 5 & 25 \\ 1 & 5 & 25 \\ 1 & 5 & 25 \\ 1 & 7 & 49 \\ 1 & 7 & 49 \\ 1 & 8 & 64 \\ \vdots & \vdots & \vdots \\ 1 & 39 & 1521 \\ 1 & 39 & 1521 \\ 1 & 40 & 1600 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 4 \\ 2 \\ 4 \\ 8 \\ 8 \\ 7 \\ 7 \\ 8 \\ \vdots \\ 138 \\ 110 \\ 134 \end{bmatrix}$$

The normal equations  $\Rightarrow \hat{\beta} = \begin{bmatrix} 1.58 \\ 0.42 \\ 0.066 \end{bmatrix}$



# Transforming the inputs

---

If the original input variable is  $v$ , we can for instance use:



- a polynomial in  $v$

$$y = \beta_0 + \beta_1 \underset{x_1}{\underset{\parallel}{v}} + \beta_2 \underset{x_2}{\underset{\parallel}{v^2}} + \beta_3 \underset{x_3}{\underset{\parallel}{v^3}} + \cdots + \beta_k \underset{x_k}{\underset{\parallel}{v^k}} + \varepsilon$$

- radial basis function kernels (see book draft)
- ...

# Ex) Happiness

Happiness is fitted as a function of Log GDP, Social Support, Life Expectancy, Freedom, Generosity and Corruption.




<https://worldhappiness.report/ed/2019/>

# Ex) Happiness

Happiness is fitted as a function of Log GDP, Social Support, Life Expectancy, Freedom, Generosity and Corruption.

Table 2.1: Regressions to Explain Average Happiness across Countries (Pooled OLS)


Independent Variable	Dependent Variable			
	Cantril Ladder (0-10)	Positive Affect (0-1)	Negative Affect (0-1)	Cantril Ladder (0-10)
Log GDP per capita	0.318 (0.066)***	-0.01 (0.01)	0.008 (0.008)	0.338 (0.065)***
Social support	2.422 (0.381)***	0.253 (0.05)***	-0.313 (0.051)***	1.977 (0.397)***
Healthy life expectancy at birth	0.033 (0.01)***	0.001 (0.001)	0.002 (0.001)	0.03 (0.01)***
Freedom to make life choices	1.164 (0.3)***	0.352 (0.04)***	-0.072 (0.041)*	0.461 (0.287)
Generosity	0.635 (0.277)**	0.137 (0.03)***	0.008 (0.028)	0.351 (0.279)
Perceptions of corruption	-0.540 (0.294)*	0.025 (0.027)	0.094 (0.024)***	-0.612 (0.287)**
Positive affect				2.063 (0.384)***
Negative affect				0.342 (0.429)
Year fixed effects	Included	Included	Included	Included
Number of countries	157	157	157	157
Number of obs.	1,516	1,513	1,515	1,512
Adjusted R-squared	0.74	0.476	0.27	0.76



<https://worldhappiness.report/ed/2019/>

# Is the model too flexible?


With a  $k = n - 1$  degree polynomial, we can fit  $n$  data points perfectly.





# Is the model too flexible?


With a  $k = n - 1$  degree polynomial, we can fit  $n$  data points perfectly.



Is this desired?

# Is the model too flexible?

With a  $k = n - 1$  degree polynomial, we can fit  $n$  data points perfectly.



Is this desired? **Overfit!**

# Regularization



*"Keep  $\beta$  small unless the data really convinces us otherwise"*

Least squares

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{X}\beta - \mathbf{y}\|_2^2$$

# Regularization



*"Keep  $\beta$  small unless the data really convinces us otherwise"*

Least squares with Ridge regression

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{X}\beta - \mathbf{y}\|_2^2 + \gamma \|\beta\|_2^2$$

# Regularization




*"Keep  $\beta$  small unless the data really convinces us otherwise"*

Least squares with Ridge regression

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{X}\beta - \mathbf{y}\|_2^2 + \gamma \|\beta\|_2^2$$
$$\Rightarrow (\mathbf{X}^\top \mathbf{X} + \gamma \mathbf{I}_{p+1}) \hat{\beta} = \mathbf{X}^\top \mathbf{y}$$

$\gamma$  regularization parameter

# Is the model too flexible?



*Regularization can help us to avoid overfitting!*

# Regularization

---

## Ridge regression

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{X}\beta - \mathbf{y}\|_2^2 + \gamma \|\beta\|_2^2$$

(has a closed-form solution for  $\hat{\beta}$ )

## LASSO

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{X}\beta - \mathbf{y}\|_2^2 + \gamma \|\beta\|_1$$

(lacks a closed-form solution for  $\hat{\beta}$ )

Regularization can be used in many methods, not only linear regression!

# Dummy variables for categorical inputs

---

For a categorical input with 2 different classes/levels/labels A and B:  
Create a dummy variable

$$x = \begin{cases} 0 & \text{if A} \\ 1 & \text{if B} \end{cases}$$

$$\Rightarrow y = \beta_0 + \beta_1 x + \varepsilon = \begin{cases} \beta_0 + \varepsilon & \text{if A} \\ \beta_0 + \beta_1 + \varepsilon & \text{if B} \end{cases}$$



# Dummy variables for categorical inputs

For a categorical input with  $a = 4$  different classes/levels/labels A, B, C, D:  
Create  $a - 1 = 3$  dummy variables

$$x_1 = \begin{cases} 1 & \text{if B} \\ 0 & \text{if not B} \end{cases}, \quad x_2 = \begin{cases} 1 & \text{if C} \\ 0 & \text{if not C} \end{cases}, \quad x_3 = \begin{cases} 1 & \text{if D} \\ 0 & \text{if not D} \end{cases}$$

$$\Rightarrow y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \varepsilon = \begin{cases} \beta_0 + \varepsilon & \text{if A} \\ \beta_0 + \beta_1 + \varepsilon & \text{if B} \\ \beta_0 + \beta_2 + \varepsilon & \text{if C} \\ \beta_0 + \beta_3 + \varepsilon & \text{if D} \end{cases}$$



# A few concepts to summarize lecture 2

---

**Regression** is about learning a model that describes the relationship between an input variable  $x$  (both numerical and categorical) and a numerical output variable  $y$ .

**Linear regression** corresponds to regression with a linear model.

**Maximum likelihood** with Gaussian iid assumption on  $\varepsilon$   
 $\Rightarrow$  **least squares** and **normal equations**

**Nonlinear transformations** can be applied to the input variables

**Overfit** is when the model adapts (too much) to noise in the data

**Regularization** can help against overfitting

**Categorical variables** are handled by dummy variables