

LogReg

February 23, 2021

0.1 Imports and data ingestion

linear_model for Logistic Regression,
model_selection for train-test split and cross-validation,
metrics for ROC and AIC

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn.linear_model as skl_lm
import sklearn.ensemble as skl_en
import sklearn.model_selection as skl_ms
import sklearn.metrics as skl_met
import sklearn.preprocessing as skl_pre
```

```
[2]: rawData = pd.read_csv('train.csv')
```

0.2 Get Dummies

Lead is 0 if the lead is female.

```
[3]: data = rawData.copy()
data['Lead'] = pd.get_dummies(data['Lead'])
```

0.3 Feature Engineering

We first extract the number of words spoken by the co-lead.

```
[4]: data['Number of words co-lead'] = data['Number of words lead'] -
    ↳data['Difference in words lead and co-lead']
```

It is probably more important to know proportion of words spoken by lead than the absolute number of words.

```
[5]: data['Proportion of words lead'] = data['Number of words lead']/data['Total_
    ↳words']
```

We look at the proportion of dialogue from co-lead as well as the ratio of dialogue between co-lead and lead.

```
[6]: data['Proportion of words co-lead'] = data['Number of words co-lead']/
      ↪data['Total words']
      data['Ratio words co-lead lead'] = data['Number of words co-lead']/data['Number_
      ↪of words lead']
```

Similarly for number of words spoken by male or female actors. However, here we normalize by Total words - Number of words lead since this only measures non-lead dialogue. We only need Proportion of words female since a hypothetical Proportion of words male is just 1 - Proportion of words female.

```
[7]: data['Proportion of words female'] = data['Number words female']/(data['Total_
      ↪words'] - data['Number of words lead'])
```

It might be important to look at the ratio of female actors.

```
[8]: data['Number of actors'] = data['Number of male actors'] + data['Number of_
      ↪female actors']
```

```
[9]: data[['Total words', 'Number of words lead', 'Difference in words lead and_
      ↪co-lead', 'Number of words co-lead', 'Number words male', 'Number words female']]
```

```
[9]:      Total words  Number of words lead  Difference in words lead and co-lead \
0          6394          2251.0          343
1          8780          2020.0         1219
2          4176           942.0          787
3          9855          3440.0         2623
4          7688          3835.0         3149
...          ...          ...          ...
1034         2398          1334.0         1166
1035         8404          1952.0          187
1036         2750           877.0          356
1037         3994           775.0           52
1038        11946          3410.0         1536
```

```
      Number of words co-lead  Number words male  Number words female
0          1908.0          2631          1512
1           801.0          5236          1524
2           155.0          3079           155
3           817.0          5342          1073
4           686.0          2536          1317
...           ...          ...          ...
1034          168.0           761           303
1035         1765.0          5820           632
1036          521.0           547          1326
1037          723.0          2757           462
1038         1874.0          5801          2735
```

```
[1039 rows x 6 columns]
```

```
[10]: data['Proportion of female actors'] = data['Number of female actors']/
      ↪data['Number of actors']
```

Perhaps it is important if the lead or co-lead is oldest.

```
[11]: data['Older lead'] = data['Age Lead'] < data['Age Co-Lead']
      data['Older lead'] = pd.get_dummies(data['Older lead'])
```

```
[12]: for col in data.columns:
      print("'" + col + "'",)
```

```
'Number words female',
'Total words',
'Number of words lead',
'Difference in words lead and co-lead',
'Number of male actors',
'Year',
'Number of female actors',
'Number words male',
'Gross',
'Mean Age Male',
'Mean Age Female',
'Age Lead',
'Age Co-Lead',
'Lead',
'Number of words co-lead',
'Proportion of words lead',
'Proportion of words co-lead',
'Ratio words co-lead lead',
'Proportion of words female',
'Number of actors',
'Proportion of female actors',
'Older lead',
```

We center and scale all computed columns as well as some of the original columns

```
[13]: cols_to_norm = [
      'Total words',
      'Year',
      'Gross',
      'Mean Age Male',
      'Mean Age Female',
      'Age Lead',
      'Age Co-Lead',
      'Number of actors'#,
      #'Proportion of words lead',
      #'Proportion of words co-lead',
      #'Ratio words co-lead lead',
      #'Proportion of words female',
```

```

    #'Proportion of female actors',
]

```

```

[14]: scaler = skl_pre.StandardScaler()
      data[cols_to_norm] = scaler.fit_transform(data[cols_to_norm])

```

```

[15]: data

```

```

[15]:
      Number words female  Total words  Number of words lead \
0          1512      -0.676591      2251.0
1          1524      -0.326435      2020.0
2           155     -1.002092       942.0
3         1073     -0.168675      3440.0
4         1317     -0.486691      3835.0
...          ...          ...          ...
1034         303     -1.263021      1334.0
1035         632     -0.381615      1952.0
1036        1326     -1.211363       877.0
1037         462     -1.028801       775.0
1038        2735      0.138188      3410.0

      Difference in words lead and co-lead  Number of male actors      Year \
0          343          2 -0.467462
1         1219          9  0.109371
2          787          7 -3.063211
3         2623         12  0.205509
4         3149          8 -1.140434
...          ...          ...          ...
1034        1166          5 -2.582517
1035         187          6 -0.755879
1036         356          2  0.013232
1037          52          8 -0.371324
1038        1536         13  0.686204

      Number of female actors  Number words male      Gross  Mean Age Male \
0              5          2631  0.203383      1.170591
1              4          5236 -0.488825     -0.413237
2              1          3079  1.746018      0.018716
3              2          5342 -0.607490     -0.912739
4              4          2536 -0.469048      0.370678
...          ...          ...          ...
1034              2          761  0.414341      0.108306
1035              2          5820  0.401156     -0.663877
1036              3          547 -0.383346     -1.901076
1037              3          2757 -0.521788      0.064425
1038              4          5801 -0.521788      0.222330

```

	...	Age Co-Lead	Lead	Number of words co-lead	\
0	...	2.451143	1	1908.0	
1	...	-0.123416	0	801.0	
2	...	0.125734	0	155.0	
3	...	-1.036970	0	817.0	
4	...	0.291835	0	686.0	
...	
1034	...	-0.953919	0	168.0	
1035	...	-0.123416	1	1765.0	
1036	...	-0.870869	0	521.0	
1037	...	-0.289517	1	723.0	
1038	...	1.039288	0	1874.0	

	Proportion of words lead	Proportion of words co-lead	\
0	0.352049	0.298405	
1	0.230068	0.091230	
2	0.225575	0.037117	
3	0.349061	0.082902	
4	0.498829	0.089230	
...	
1034	0.556297	0.070058	
1035	0.232270	0.210019	
1036	0.318909	0.189455	
1037	0.194041	0.181022	
1038	0.285451	0.156873	

	Ratio words co-lead lead	Proportion of words female	Number of actors	\
0	0.847623	0.364953	-0.940919	
1	0.396535	0.225444	0.379885	
2	0.164544	0.047928	-0.720785	
3	0.237500	0.167264	0.600019	
4	0.178879	0.341812	0.159751	
...	
1034	0.125937	0.284774	-0.940919	
1035	0.904201	0.097954	-0.720785	
1036	0.594071	0.707955	-1.381187	
1037	0.932903	0.143523	-0.060383	
1038	0.549560	0.320408	1.260421	

	Proportion of female actors	Older lead
0	0.714286	0
1	0.307692	1
2	0.125000	1
3	0.142857	1
4	0.333333	0
...
1034	0.285714	1

1035	0.250000	0
1036	0.600000	1
1037	0.272727	0
1038	0.235294	0

[1039 rows x 22 columns]

Combining all pre-processing into a function

```
[16]: def pre_process(raw_data, cols_to_norm):
    data = raw_data.copy()

    data['Lead'] = pd.get_dummies(data['Lead'])
    data['Number of words co-lead'] = data['Number of words lead'] -
    data['Difference in words lead and co-lead']
    data['Proportion of words lead'] = data['Number of words lead']/data['Total
    words']
    data['Proportion of words co-lead'] = data['Number of words co-lead']/
    data['Total words']
    data['Ratio words co-lead lead'] = data['Number of words co-lead']/
    data['Number of words lead']
    data['Proportion of words female'] = data['Number words female']/
    (data['Total words'] - data['Number of words lead'])
    data['Number of actors'] = data['Number of male actors'] + data['Number of
    female actors']
    data['Proportion of female actors'] = data['Number of female actors']/
    data['Number of actors']
    data['Older lead'] = data['Age Lead'] < data['Age Co-Lead']
    data['Older lead'] = pd.get_dummies(data['Older lead'])

    scaler = skl_pre.StandardScaler()
    data[cols_to_norm] = scaler.fit_transform(data[cols_to_norm])

    return data
```

1 Logistic Regression

```
[17]: data = pre_process(rawData, cols_to_norm)
```

```
[205]: config = pd.read_csv('config.csv')
trainRatio = config['Train Ratio'][0]
seed = config['Random Seed'][0]
train, test = skl_ms.train_test_split(data, train_size=trainRatio) #,
    random_state=seed)
```

```
[206]: rawFeatures = [
    'Number words female',
    'Total words',
    'Number of words lead',
    'Difference in words lead and co-lead',
    'Number of male actors',
    'Year',
    'Number of female actors',
    'Number words male',
    'Gross',
    'Mean Age Male',
    'Mean Age Female',
    'Age Lead',
    'Age Co-Lead'
]

featureSet1 = [
    'Year',
    'Gross',
    'Number of actors',
    'Proportion of female actors',
    'Mean Age Male',
    'Mean Age Female',
    'Age Lead',
    'Age Co-Lead',
    'Total words',
    'Proportion of words lead',
    'Proportion of words co-lead',
    'Ratio words co-lead lead',
    'Proportion of words female',
    'Older lead'
]
```

```
[207]: set(featureSet1) - set(cols_to_norm)
```

```
[207]: {'Older lead',
    'Proportion of female actors',
    'Proportion of words co-lead',
    'Proportion of words female',
    'Proportion of words lead',
    'Ratio words co-lead lead'}
```

```
[208]: features = featureSet1.copy()
#features.remove('Proportion of words female')
#features.remove('Year')
#features.remove('Gross')
#features = ['Proportion of words lead']
```

```
target = 'Lead'
```

```
[209]: def fit_and_test(classifier, train, test, features, target, suppress_output =  
→False):  
    classifier.fit(train[features], train[target])  
    if not suppress_output:  
        skl_met.plot_roc_curve(classifier, test[features], test[target])  
        print('accuracy: ' + str(classifier.score(test[features], test[target])))  
        print('    auc: ' + str(skl_met.roc_auc_score(test[target], classifier.  
→predict_proba(test[features]))[:,1])) + '\n')  
        print(skl_met.classification_report(test[target], classifier.  
→predict(test[features])))  
    return classifier
```

```
[210]: print('Null accuracy: ' + str(max([np.mean(data[target]), 1 - np.  
→mean(data[target])])))
```

Null accuracy: 0.7555341674687199

1.1 No regularization

```
[211]: %%time  
B = 100  
accuracies = []  
aucs = []  
for i in range(B):  
    train, test = skl_ms.train_test_split(data, train_size=trainRatio)  
    logReg = fit_and_test(skl_lm.LogisticRegression(penalty='none',  
→solver='newton-cg'), train, test, features, target, suppress_output=True)  
    accuracies.append(logReg.score(test[features], test[target]))  
    aucs.append(skl_met.roc_auc_score(test[target], logReg.  
→predict_proba(test[features]))[:,1]))
```

CPU times: user 8.2 s, sys: 21.2 s, total: 29.4 s

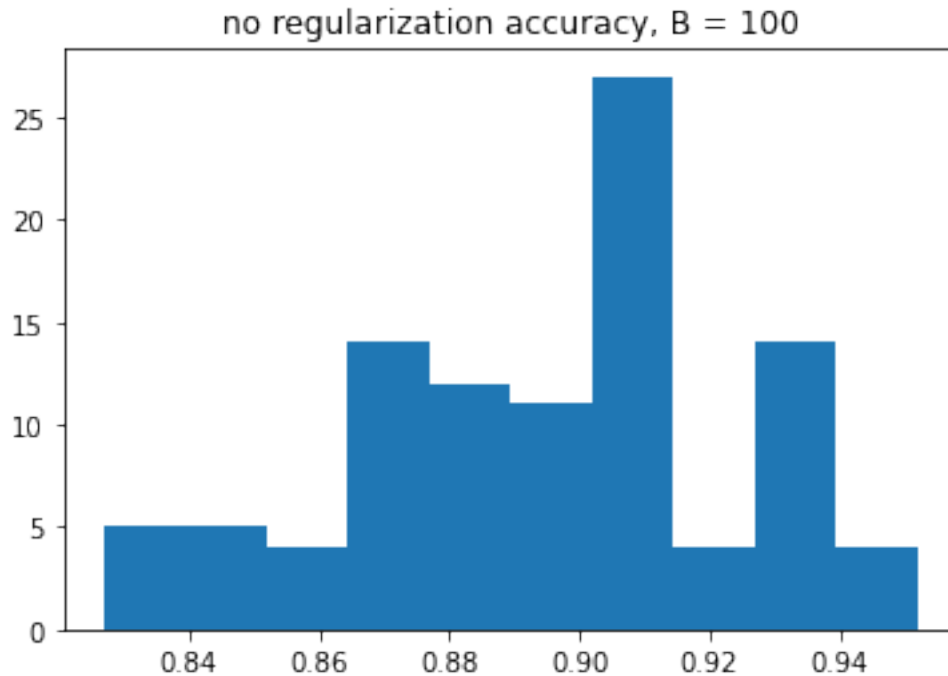
Wall time: 2.6 s

```
[212]: print('mean accuracy: ' + str(np.mean(accuracies)))  
print('    mean auc: ' + str(np.mean(aucs)))
```

mean accuracy: 0.8953846153846154

mean auc: 0.9244967487778538

```
[213]: plt.hist(accuracies)  
plt.title(f'no regularization accuracy, B = {B}')  
plt.show()
```

1.2 Lasso

Lasso is L1 regularization

We do 20 folds, use `liblinear` solver that is better for smaller datasets and 10 threads in the cross-validation step.


The method will try 10 values of `lambda` between $1e-4$ and $1e4$ to determine the best regularization parameter.

```
[214]: logRegLasso = fit_and_test(skl_lm.LogisticRegressionCV(Cs=10, cv=50,
    ↳penalty='l1', solver='liblinear', n_jobs=10), train, test, features, target)
```

accuracy: 0.9038461538461539

auc: 0.9235033259423504

	precision	recall	f1-score	support
0	0.92	0.96	0.94	82
1	0.83	0.68	0.75	22
accuracy			0.90	104
macro avg	0.88	0.82	0.85	104
weighted avg	0.90	0.90	0.90	104



We want to know how the algorithm performs over different splits of training and testing data

```
[215]: %%time
B = 100
accuracies = []
aucs = []
for i in range(B):
    train, test = skl_ms.train_test_split(data, train_size=trainRatio)
    logRegLasso = fit_and_test(skl_lm.LogisticRegressionCV(Cs=10, cv=10,
    →penalty='l1', solver='liblinear', n_jobs=10), train, test, features, target,
    →suppress_output=True)
    accuracies.append(logRegLasso.score(test[features], test[target]))
    aucs.append(skl_met.roc_auc_score(test[target], logRegLasso.
    →predict_proba(test[features]))[:,1]))
```

CPU times: user 4.31 s, sys: 46.4 ms, total: 4.36 s

Wall time: 23.2 s

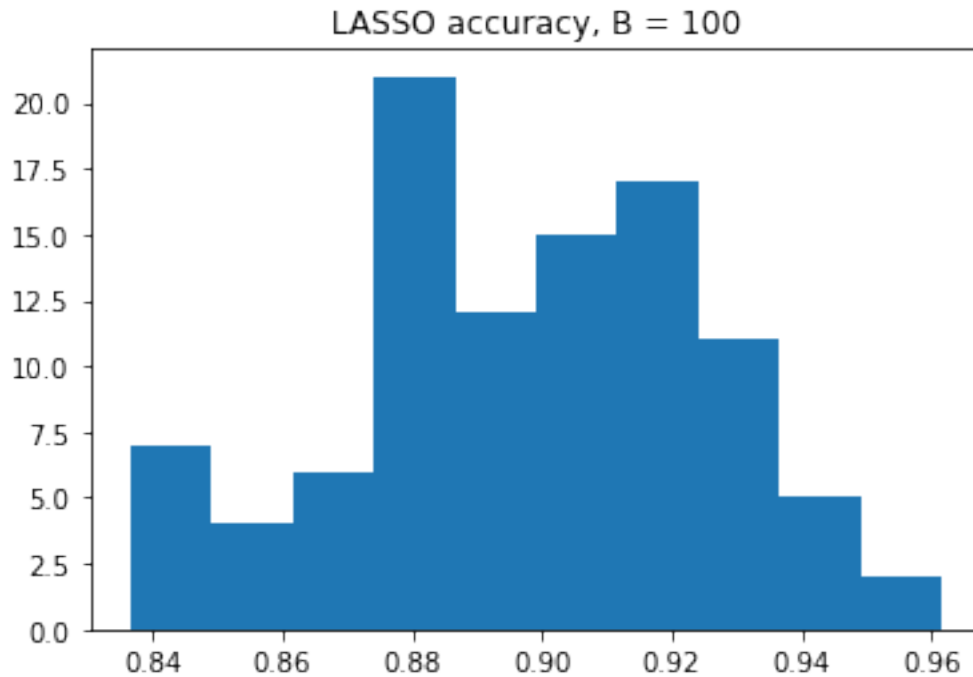
```
[216]: print('mean accuracy: ' + str(np.mean(accuracies)))
print('      mean auc: ' + str(np.mean(aucs)))
```

mean accuracy: 0.8972115384615386

mean auc: 0.9244666002902575

```
[217]: plt.hist(accuracies)
plt.title(f'LASSO accuracy, B = {B}')
```

```
plt.show()
```



1.3 Ridge

Ridge is L2-regularization

We do 10 folds, use `liblinear` solver that is better for smaller datasets and 10 threads in the cross-validation step.

The method will try 100 values of `lambda` between $1e-4$ and $1e4$ to determine the best regularization parameter.


```
[218]: logRegRidge = fit_and_test(skl_lm.LogisticRegressionCV(Cs=10, cv=20,  
    ↪penalty='l2', solver='liblinear', n_jobs=10), train, test, features, target)
```

```
accuracy: 0.9038461538461539
```

```
auc: 0.9098734177215191
```

	precision	recall	f1-score	support
0	0.93	0.95	0.94	79
1	0.83	0.76	0.79	25
accuracy			0.90	104
macro avg	0.88	0.85	0.86	104

weighted avg 0.90 0.90 0.90 104



```
[219]: %%time
B = 400
accuracies = []
aucs = []
for i in range(B):
    train, test = skl_ms.train_test_split(data, train_size=trainRatio)
    logRegLasso = fit_and_test(skl_lm.LogisticRegressionCV(Cs=10, cv=10,
    →penalty='l2', solver='liblinear', n_jobs=10), train, test, features, target,
    →suppress_output=True)
    accuracies.append(logRegLasso.score(test[features], test[target]))
    aucs.append(skl_met.roc_auc_score(test[target], logRegLasso.
    →predict_proba(test[features]))[:,1]))
```

CPU times: user 5.96 s, sys: 233 ms, total: 6.19 s

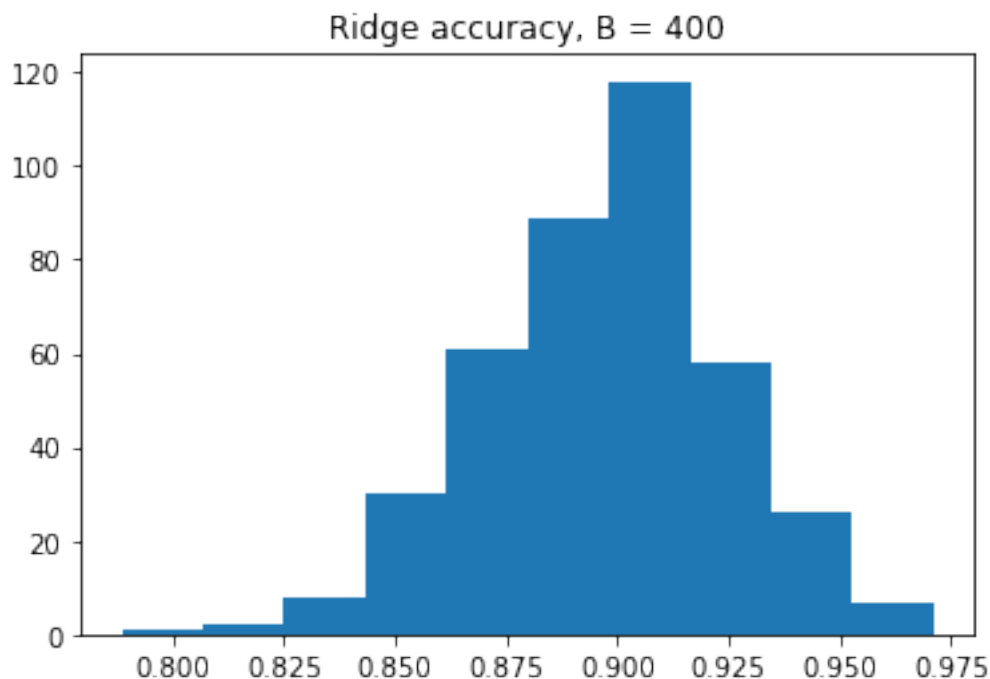
Wall time: 20 s

```
[220]: print("mean accuracy: " + str(np.mean(accuracies)))
print('    mean auc: ' + str(np.mean(aucs)))
```

mean accuracy: 0.8977884615384616

mean auc: 0.9231520433528324

```
[221]: plt.hist(accuracies)
plt.title(f'Ridge accuracy, B = {B}')
plt.show()
```



1.4 Comparison

```
[222]: logRegLasso.C_
```

```
[222]: array([10000.])
```

```
[223]: logRegRidge.C_
```

```
[223]: array([21.5443469])
```

```
[224]: predLasso = logRegLasso.predict(test[features])
predRidge = logRegRidge.predict(test[features])
```

```
[225]: logRegLassoCoefs = dict(zip(features, logRegLasso.coef_[0]))
```

There is a reason that Proportion of words female is an indicator that the lead is male. If the Lead is male, Number words female includes the number of words spoken by the co-lead and hence Number words female is the total number of words spoken by **all** female actors while Number words male is the number of words spoken by the **non-lead** male actors. Hence,

```
[226]: for var, coef in logRegLassoCoefs.items():
        print(f'{var:40} {coef:.5f}')
```

```
Year                                -0.07222
Gross                              -0.02181
Number of actors                    0.56105
Proportion of female actors        18.62474
Mean Age Male                       0.19473
Mean Age Female                     0.05373
Age Lead                           -0.30258
Age Co-Lead                         0.21755
Total words                         -0.58851
Proportion of words lead            2.64110
Proportion of words co-lead         3.81626
Ratio words co-lead lead            3.30966
Proportion of words female          -11.55700
Older lead                          -1.21839
```

```
[227]: logRegRidgeCoefs = dict(zip(features, logRegRidge.coef_[0]))
```

```
[228]: for var, coef in logRegRidgeCoefs.items():
        print(f'{var:40} {coef:.5f}')
```

```
Year                                -0.05445
Gross                              -0.04454
Number of actors                    0.41936
Proportion of female actors        16.99463
Mean Age Male                       0.15562
Mean Age Female                     0.09000
Age Lead                           -0.32796
Age Co-Lead                         0.19087
Total words                         -0.48615
Proportion of words lead            2.14504
Proportion of words co-lead         4.16098
Ratio words co-lead lead            2.88327
Proportion of words female          -10.79023
Older lead                          -1.23370
```

```
[229]: sum(predLasso == predRidge)/len(test)
```

```
[229]: 0.9615384615384616
```

```
[230]: pd.crosstab(test[target], predLasso)
```

```
[230]: col_0    0    1
Lead
0         69    4
1          5   26
```

```
[231]: pd.crosstab(test[target], predRidge)
```

```
[231]: col_0    0    1  
Lead  
0       71    2  
1        5   26
```

```
[ ]:
```